

Self-Driving Car Engineer Nanodegree

Deep Learning

Project: Build a Traffic Sign Recognition Classifier

This Jupyter notebook contains the code pipeline for a traffic sign classifier tuned for project 2 of Term 1. Python libraries used in this project:

- `pickle`: save and load binary python objects
- `numpy`: algebra calculations
- `matplotlib`: plots and image loading
- `tensorflow`: machine learning framework
- `sklearn`: machine learning framework

Packages `scikit-image` and `cv2` were tested, but not used in final form.

Step 0: Load The Data

Input data was provided in binary files.

```
In [1]: # Load pickled data
import pickle
import numpy as np

training_file = './traffic-signs-data/train.p'
validation_file = './traffic-signs-data/valid.p'
testing_file = './traffic-signs-data/test.p'

with open(training_file, mode='rb') as f:
    train = pickle.load(f)
with open(validation_file, mode='rb') as f:
    valid = pickle.load(f)
with open(testing_file, mode='rb') as f:
    test = pickle.load(f)

X_train0, y_train0 = train['features'], train['labels']
X_valid0, y_valid0 = valid['features'], valid['labels']
X_test0, y_test0 = test['features'], test['labels']
```

Step 1: Dataset Summary & Exploration

The pickled data is a dictionary with 4 key/value pairs:

- 'features' is a 4D array containing raw pixel data of the traffic sign images, (num examples, width, height, channels).
- 'labels' is a 1D array containing the label/class id of the traffic sign. The file `signnames.csv` contains id -> name mappings for each id.

Basic Summary of the Data Set Using Python and Numpy

```
In [2]: # Number of training examples
n_train = X_train0.shape[0]

# Number of validation examples
n_validation = X_valid0.shape[0]

# Number of testing examples.
n_test = X_test0.shape[0]

# What's the shape of an traffic sign image?
image_shape = X_train0.shape[1], X_train0.shape[2], X_train0.shape[3]

# How many unique classes/labels there are in the dataset.
n_classes = np.max(np.unique(y_train0).shape)

print("Number of training examples =", n_train)
print("Number of validation examples =", n_validation)
print("Number of testing examples =", n_test)
print("Image data shape =", image_shape)
print("Number of classes =", n_classes)
print("labels shape: ", y_train0.shape)
print(y_train0[0:4])
```

```
Number of training examples = 34799
Number of validation examples = 4410
Number of testing examples = 12630
Image data shape = (32, 32, 3)
Number of classes = 43
labels shape: (34799,)
[41 41 41 41]
```

Visualization of the dataset

Next, it is presented 4 random images from the training dataset and a histogram of labels (output).

```
In [3]: ### Data exploration visualization code goes here.
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
# Visualizations will be shown in the notebook.
%matplotlib inline

im1 = np.round(np.random.random_sample()*n_train).astype(int)
im2 = np.round(np.random.random_sample()*n_train).astype(int)
im3 = np.round(np.random.random_sample()*n_train).astype(int)
im4 = np.round(np.random.random_sample()*n_train).astype(int)


print(im1, im2, im3, im4)

fig, axs = plt.subplots(nrows=2, ncols=2)
axs[0, 0].imshow(X_train0[im1,:,:,:])
axs[0, 1].imshow(X_train0[im2,:,:,:])
axs[1, 0].imshow(X_train0[im3,:,:,:])
axs[1, 1].imshow(X_train0[im4,:,:,:])

plt.show()

plt.figure()
plt.hist(train['labels'], n_classes)
```


20072 9319 18600 22108



```

Out[3]: (array([ 180., 1980., 2010., 1260., 1770., 1650., 360., 1
290.,
          1260., 1320., 1800., 1170., 1890., 1920., 690.,
540.,
          360., 990., 1080., 180., 300., 270., 330.,
450.,
          240., 1350., 540., 210., 480., 240., 390.,
690.,
          210., 599., 360., 1080., 330., 180., 1860.,
270.,
          300., 210., 210.]),
array([ 0.,          0.97674419, 1.95348837, 2.93023256,
3.90697674, 4.88372093, 5.86046512, 6.8372093 ,
7.81395349, 8.79069767, 9.76744186, 10.74418605,
11.72093023, 12.69767442, 13.6744186 , 14.65116279,
15.62790698, 16.60465116, 17.58139535, 18.55813953,
19.53488372, 20.51162791, 21.48837209, 22.46511628,
23.44186047, 24.41860465, 25.39534884, 26.37209302,
27.34883721, 28.3255814 , 29.30232558, 30.27906977,
31.25581395, 32.23255814, 33.20930233, 34.18604651,
35.1627907 , 36.13953488, 37.11627907, 38.09302326,
39.06976744, 40.04651163, 41.02325581, 42.          ]),
<a list of 43 Patch objects>)

```



Labels in the training set are not equally distributed.

Step 2: Design and Test a Model Architecture

This project is based on LeNet-5 architecture. LeNet consists of 2 convolutional layers, 1 max pool and 2 fully connected layer.

The architecture proposed here adds another fully connected layer and uses a deeper convolutional filter.

Pre-process the Data Set

Pre-processing is made in two steps:

- grayscaling with: $\text{Gray} = 0.299 \text{ Red} + 0.587 \text{ Green} + 0.114 \text{ Blue}$
- normalizing pixel intensity

```
In [4]: ### Preprocess the data here.
from sklearn.utils import shuffle
#from skimage import exposure
#from skimage.color import rgb2gray
from datetime import datetime

# training data
X_train = np.zeros((n_train, X_train0.shape[1], X_train0.shape[2],
1), np.int)

# validation data
X_valid = np.zeros((n_validation, X_valid0.shape[1], X_valid0.shape
[2], 1), np.int)

# grayscale luminosity
#lum = np.ndarray((3,), np.float, np.array([0.210, 0.720, 0.070]))
lum = np.ndarray((3,), np.float, np.array([0.299, 0.587, 0.114]))

y_train = y_train0
y_valid = y_valid0

for i in range(n_train):
    # grayscale with CV2
    #X_train[i,:,:,:] = cv2.cvtColor(X_train[i,:,:,:], cv2.COLOR_RGB2GRAY)
    # grayscale with luminosity
    # 0.21 R + 0.72 G + 0.07 B.
    #X_train[i,:,:,:] = (X_train1[i,:,:,:] * 0.21 + X_train1[i,:,:,:] * 0.72 + X_train1[i,:,:,:] * 0.07).astype(int)
    # also suggested in CarND online forums
    #X_train[i,:,:,:] = np.dot(X_train1[i,:,:,:], lum).astype(int)
```

```

#p2, p98 = np.percentile(X_train1[i,:,:,:], (2, 98))
#X_train1[i,:,:,:] = exposure.rescale_intensity(X_train1[i,:,:,:],
:], in_range=(p2, p98))
#pass
#X_train[i,:,:,:0] = np.dot(X_train1[i,:,:,:],
#X_train[i,:,:,:0] = rgb2gray(X_train1[i,:,:,:])

# grayscale
X_train[i,:,:,:0] = np.dot(X_train0[i,:,:,:], lum)

for i in range(n_validation):
    X_valid[i,:,:,:0] = np.dot(X_valid0[i,:,:,:], lum).astype(int)

# normalized after grayscale and save computing costs
X_train = (X_train/255.0)-0.5
X_valid = (X_valid/255.0)-0.5

# check data
print("max min ", np.amin(X_train), np.amax(X_train))


fig, axs = plt.subplots(nrows=2, ncols=2)
axs[0, 0].imshow(X_train[im1,:,:,:0], cmap=plt.cm.gray)
axs[0, 1].imshow(X_train[im2,:,:,:0], cmap=plt.cm.gray)
axs[1, 0].imshow(X_train[im3,:,:,:0], cmap=plt.cm.gray)
axs[1, 1].imshow(X_train[im4,:,:,:0], cmap=plt.cm.gray)

print("Normalized grayscaled images.")
plt.show()

print("Done preprocessing.")

```

max min -0.488235294118 0.5
Normalized grayscaled images.



Done preprocessing.

Model Architecture

The python function which creates the logitis was coded with a factory pattern, a form of closure. This makes it reusable with different number of output classes.

```
In [5]: import tensorflow as tf
from tensorflow.contrib.layers import flatten
from tensorflow.python.client import device_lib

import platform

def factory(n_classes, mu = 0, sigma = 0.1):

    def LeNet(x):

        # W=32, F=5, P=0, S=1
        # out = 1 + [W-F+2P]/S => 1 + (32-5+0)/1 = 28
        # Input = 32x32x1. Output = 28x28x6.
        # number of filters is arbitrary
        # https://discussions.udacity.com/t/define-input-depth-output-depth-f/238575/14
        #
        conv1_W = tf.Variable(tf.truncated_normal(shape=(5, 5, 1, 16), mean = mu, stddev = sigma))
        conv1_b = tf.Variable(tf.zeros(16))
        conv1 = tf.nn.conv2d(x, conv1_W, strides=[1, 1, 1, 1], padding='VALID') + conv1_b

        # Activation.
        conv1 = tf.nn.relu(conv1)

        # out = 1 + [W-F+2P]/S => 1+(28-2+0)/2 = 14
        # Pooling. Input = 28x28x16, Output = 14x14x16.
        conv1 = tf.nn.max_pool(conv1, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='VALID')

        # out = 1 + [W-F+2P]/S => 1+(14-5+0)/1 = 10
        # Layer 2: Convolutional. Input = 14x14x16, Output = 10x10x32.
        conv2_W = tf.Variable(tf.truncated_normal(shape=(5, 5, 16, 32), mean = mu, stddev = sigma))
        conv2_b = tf.Variable(tf.zeros(32))
        conv2 = tf.nn.conv2d(conv1, conv2_W, strides=[1, 1, 1, 1], padding='VALID') + conv2_b

        # Activation.
        conv2 = tf.nn.relu(conv2)

        # out = 1 + [W-F+2P]/S => 1+(10-2+0)/2 = 5
        # Pooling. Input = 10x10x16, Output = 5x5x32.
```

```

        conv2 = tf.nn.max_pool(conv2, ksize=[1, 2, 2, 1], strides=[
1, 2, 2, 1], padding='VALID')

        # Flatten. Input = 5x5x32. Output = 800.
        fc0    = flatten(conv2)

        # out = 1 + [W-F+2P]/S =>
        # Fully Connected. Input = 800. Output = 400.
        fc1_W = tf.Variable(tf.truncated_normal(shape=(800, 400), m
ean = mu, stddev = sigma))
        fc1_b = tf.Variable(tf.zeros(400))
        fc1    = tf.matmul(fc0, fc1_W) + fc1_b

        # Activation.
        fc1     = tf.nn.relu(fc1)

        # connected layer
        fc2_W = tf.Variable(tf.truncated_normal(shape=(400, 129), m
ean = mu, stddev = sigma))
        fc2_b = tf.Variable(tf.zeros(129))
        fc2    = tf.matmul(fc1, fc2_W) + fc2_b

        fc2     = tf.nn.relu(fc2)

        # Fully Connected. Input = 129. Output = 86.
        fc3_W = tf.Variable(tf.truncated_normal(shape=(129, 86), m
ean = mu, stddev = sigma))
        fc3_b = tf.Variable(tf.zeros(86))
        fc3    = tf.matmul(fc2, fc3_W) + fc3_b

        # Activation.
        fc3     = tf.nn.relu(fc3)

        # Input = 86 Output = n_classes.
        fc4_W = tf.Variable(tf.truncated_normal(shape=(86, n_class
es), mean = mu, stddev = sigma))
        fc4_b = tf.Variable(tf.zeros(n_classes))

        # final
        logits = tf.matmul(fc3, fc4_W) + fc4_b

    return logits
return LeNet

```

Train, Validate and Test the Model

Input data was already split into training, validation and testing. This separation helps to prevent overfitting.

Running the model on AWS enables computations on GPU. For 50 epochs it takes about 3 minutes.

```
In [6]: def evaluate(X_data, y_data):
        num_examples = len(X_data)
        total_accuracy = 0
        total_loss = 0
        sess = tf.get_default_session()
        for offset in range(0, num_examples, BATCH_SIZE):
            batch_x, batch_y = X_data[offset:offset+BATCH_SIZE], y_data[
offset:offset+BATCH_SIZE]
            loss, accuracy = sess.run([loss_operation, accuracy_operati
on], feed_dict={x: batch_x, y: batch_y})
            total_loss += (loss*len(batch_x))
            total_accuracy += (accuracy * len(batch_x))
        return total_loss/num_examples, total_accuracy / num_examples
```

```

In [7]: ### Training pipeline

x = tf.placeholder(tf.float32, (None, 32, 32, 1))
y = tf.placeholder(tf.int32, (None))
one_hot_y = tf.one_hot(y, n_classes)

EPOCHS = 50
BATCH_SIZE = 128
#BATCH_SIZE = 256
rate = 0.0005

LeNetFn = factory(n_classes)
logits = LeNetFn(x)
cross_entropy = tf.nn.softmax_cross_entropy_with_logits(labels=one_hot_y, logits=logits)
loss_operation = tf.reduce_mean(cross_entropy)
optimizer = tf.train.AdamOptimizer(learning_rate = rate)
training_operation = optimizer.minimize(loss_operation)

correct_prediction = tf.equal(tf.argmax(logits, 1), tf.argmax(one_hot_y, 1))
accuracy_operation = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
saver = tf.train.Saver()

print("System: ")
print(platform.uname())
print("")
devices = [x.name for x in device_lib.list_local_devices() if x.device_type == 'GPU']
print(devices)
print("")

acc_epochs = np.zeros((EPOCHS,), np.float)
loss_fn = np.zeros((EPOCHS,), np.float)

System:
uname_result(system='Linux', node='ip-172-31-32-67', release='4.4.0-97-generic', version='#120-Ubuntu SMP Tue Sep 19 17:28:18 UTC 2017', machine='x86_64', processor='x86_64')

['/gpu:0']

```

```
In [8]: with tf.Session() as sess:
        sess.run(tf.global_variables_initializer())
        num_examples = len(X_train)

        print(datetime.now().isoformat(' '), " - Training...")
        print()
        for i in range(EPOCHS):
            X_train, y_train = shuffle(X_train, y_train)
            for offset in range(0, num_examples, BATCH_SIZE):
                end = offset + BATCH_SIZE
                batch_x, batch_y = X_train[offset:end], y_train[offset:
end]
                sess.run(training_operation, feed_dict={x: batch_x, y:
batch_y})

                loss, validation_accuracy = evaluate(X_valid, y_valid)
                print("EPOCH {} ...".format(i+1))
                print("Validation Accuracy = {:.3f}".format(validation_accu
racy))
                print()
                acc_epochs[i] = validation_accuracy
                loss_fn[i] = loss

        print(datetime.now().isoformat(' '), " - Finished training")
        saver.save(sess, './lenet/lenet')
        print("Model saved")
```

2017-10-12 14:08:35.909454 - Training...

EPOCH 1 ...
Validation Accuracy = 0.703

EPOCH 2 ...
Validation Accuracy = 0.847

EPOCH 3 ...
Validation Accuracy = 0.867

EPOCH 4 ...
Validation Accuracy = 0.892

EPOCH 5 ...
Validation Accuracy = 0.890

EPOCH 6 ...
Validation Accuracy = 0.904

EPOCH 7 ...
Validation Accuracy = 0.911

EPOCH 8 ...
Validation Accuracy = 0.914

```
EPOCH 9 ...  
Validation Accuracy = 0.903  
  
EPOCH 10 ...  
Validation Accuracy = 0.909  
  
EPOCH 11 ...  
Validation Accuracy = 0.915  
  
EPOCH 12 ...  
Validation Accuracy = 0.902  
  
EPOCH 13 ...  
Validation Accuracy = 0.900  
  
EPOCH 14 ...  
Validation Accuracy = 0.914  
  
EPOCH 15 ...  
Validation Accuracy = 0.906  
  
EPOCH 16 ...  
Validation Accuracy = 0.917  
  
EPOCH 17 ...  
Validation Accuracy = 0.925  
  
EPOCH 18 ...  
Validation Accuracy = 0.908  
  
EPOCH 19 ...  
Validation Accuracy = 0.931  
  
EPOCH 20 ...  
Validation Accuracy = 0.932  
  
EPOCH 21 ...  
Validation Accuracy = 0.927  
  
EPOCH 22 ...  
Validation Accuracy = 0.922  
  
EPOCH 23 ...  
Validation Accuracy = 0.933  
  
EPOCH 24 ...  
Validation Accuracy = 0.929  
  
EPOCH 25 ...  
Validation Accuracy = 0.920  
  
EPOCH 26 ...  
Validation Accuracy = 0.908
```

```
EPOCH 27 ...  
Validation Accuracy = 0.925  
  
EPOCH 28 ...  
Validation Accuracy = 0.914  
  
EPOCH 29 ...  
Validation Accuracy = 0.935  
  
EPOCH 30 ...  
Validation Accuracy = 0.939  
  
EPOCH 31 ...  
Validation Accuracy = 0.931  
  
EPOCH 32 ...  
Validation Accuracy = 0.944  
  
EPOCH 33 ...  
Validation Accuracy = 0.944  
  
EPOCH 34 ...  
Validation Accuracy = 0.940  
  
EPOCH 35 ...  
Validation Accuracy = 0.943  
  
EPOCH 36 ...  
Validation Accuracy = 0.942  
  
EPOCH 37 ...  
Validation Accuracy = 0.940  
  
EPOCH 38 ...  
Validation Accuracy = 0.941  
  
EPOCH 39 ...  
Validation Accuracy = 0.941  
  
EPOCH 40 ...  
Validation Accuracy = 0.941  
  
EPOCH 41 ...  
Validation Accuracy = 0.942  
  
EPOCH 42 ...  
Validation Accuracy = 0.940  
  
EPOCH 43 ...  
Validation Accuracy = 0.940  
  
EPOCH 44 ...  
Validation Accuracy = 0.938
```

```
EPOCH 45 ...  
Validation Accuracy = 0.938
```

```
EPOCH 46 ...  
Validation Accuracy = 0.894
```

```
EPOCH 47 ...  
Validation Accuracy = 0.918
```

```
EPOCH 48 ...  
Validation Accuracy = 0.934
```


```
EPOCH 49 ...  
Validation Accuracy = 0.941
```

```
EPOCH 50 ...  
Validation Accuracy = 0.937
```


```
2017-10-12 14:11:29.978169 - Finished training  
Model saved
```

```
In [9]: model_r_max = np.max(acc_epochs)  
print("Average accuracy: ", np.mean(acc_epochs), " highest acc: ",  
      model_r_max)  
plt.figure()  
plt.plot(acc_epochs)  
plt.xlabel("epochs")  
plt.ylabel("accuracy")  
plt.title("Model Accuracy over Epochs")  
plt.show()  
plt.savefig("./output/evolution_1FD32_L5_E"+str(EPOCHS)+"_B"+str(BA  
TCH_SIZE)+"_R"+str(rate)+"_A94.png")  
  
plt.figure()  
plt.plot(loss_fn)  
plt.xlabel("epochs")  
plt.ylabel("loss function")  
plt.show()  
plt.savefig("./output/loss_1FD32_L5_E"+str(EPOCHS)+"_B"+str(BATCH_S  
IZE)+"_R"+str(rate)+"_A94.png")
```

Average accuracy: 0.917446711906 highest acc: 0.943764172092



<matplotlib.figure.Figure at 0x7f2fdbd10470>



<matplotlib.figure.Figure at 0x7f2fc235ef98>

Evaluate

After training and validating, its time to test the model with "new" data. This data set was provided in a separated file, so the model was not exposed to it yet.

```
In [10]: n_test = X_test0.shape[0]
X_test = np.zeros((n_test, X_test0.shape[1], X_test0.shape[2], 1),
np.int)

for i in range(n_test):
    X_test[i,:,:,:0] = np.dot(X_test0[i,:,:,:], lum).astype(int)

# normalized after grayscale and save computing costs
X_test = (X_test/255.0)-0.5

with tf.Session() as sess:
    saver.restore(sess, tf.train.latest_checkpoint('./lenet'))

    _, test_accuracy = evaluate(X_test, y_test0)
    print("Test Accuracy = {:.3f}".format(test_accuracy))
```

```
INFO:tensorflow:Restoring parameters from ./lenet/lenet
Test Accuracy = 0.935
```

Step 3: Test a Model on New Images

Ater testing the model, let's try its accuracy on 5 new images found on internet.

Load and Output the Images


```

In [11]: import matplotlib.image as mpimg

xtest20 = np.ndarray((5, 32, 32, 3), np.float)
xtest21 = np.ndarray((5, 32, 32, 1), np.float)
#xtest20[0, :, :, :] = mpimg.imread('./new-data/ni_01_14.png')
#xtest20[1, :, :, :] = mpimg.imread('./new-data/ni_02_15.png')
#xtest20[2, :, :, :] = mpimg.imread('./new-data/ni_03_27.png')
#xtest20[3, :, :, :] = mpimg.imread('./new-data/ni_04_30.png')
#xtest20[4, :, :, :] = mpimg.imread('./new-data/ni_05_40.png')

xtest20[0, :, :, :] = mpimg.imread('./new-data/ni_01_14.jpg')
xtest20[1, :, :, :] = mpimg.imread('./new-data/ni_02_15.jpg')
xtest20[2, :, :, :] = mpimg.imread('./new-data/ni_03_27.jpg')
xtest20[3, :, :, :] = mpimg.imread('./new-data/ni_04_30.jpg')
xtest20[4, :, :, :] = mpimg.imread('./new-data/ni_05_40.jpg')

ylabels = np.ndarray((5,), np.int, np.array([14,15,27,30,40]))
lum = np.ndarray((3,), np.float, np.array([0.299, 0.587, 0.114]))
for i in range(5):
    xtest21[i, :, :, 0] = np.dot(xtest20[i, :, :, :], lum)

# normalized after grayscale and save computing costs
xtest21 = (xtest21/255.0)-0.5


fig, axs = plt.subplots(nrows=3, ncols=2)
axs[0, 0].imshow(xtest20[0, :, :, :])
axs[0, 1].imshow(xtest20[1, :, :, :])
axs[1, 0].imshow(xtest20[2, :, :, :])
axs[1, 1].imshow(xtest20[3, :, :, :])
axs[2, 0].imshow(xtest20[4, :, :, :])

plt.show()

fig, axs = plt.subplots(nrows=3, ncols=2)
axs[0, 0].imshow(xtest21[0, :, :, 0], cmap=plt.cm.gray)
axs[0, 1].imshow(xtest21[1, :, :, 0], cmap=plt.cm.gray)
axs[1, 0].imshow(xtest21[2, :, :, 0], cmap=plt.cm.gray)
axs[1, 1].imshow(xtest21[3, :, :, 0], cmap=plt.cm.gray)
axs[2, 0].imshow(xtest21[4, :, :, 0], cmap=plt.cm.gray)

plt.show()

```



Predict the Sign Type for Each Image

```
In [16]: with tf.Session() as sess:
    saver.restore(sess, tf.train.latest_checkpoint('./lenet'))
    r1 = sess.run(logits, feed_dict={x: xtest21})
    print("Logits")
    print(r1)
    print(r1.shape)

    print("")
    print("")
    print("softmax")
    r2 = sess.run(tf.nn.softmax(logits), feed_dict={x: xtest21})
    print(r2)
    print(r2.shape)
    for i in range(5):
        print(i, ylabels[i], r2[i, ylabels[i]])
```

INFO:tensorflow:Restoring parameters from ./lenet/lenet

Logits

```

[[ -23.82995796    12.22348881     5.24700212    -5.01921177   -12.051
56803
   -16.28949165   -26.54062653   -19.41469574   -15.23647499   -28.703
47595
   -30.18368912    -9.78508854     6.68434906     8.00753307    14.684
8774
   -13.73761368   -65.58116913   -11.43909836   -20.74489212   -43.697
57843
   -25.68615532   -35.35309982   -25.01387787   -42.83092499   -18.187
36267
   -9.23015785    -6.61341381   -49.62615204   -24.12494278   -22.123
5199
   -27.21355057   -33.24899292   -24.66412926     4.25927591   -11.667
42802
   -2.87310266   -10.11999893    -0.73470086   -10.88895035    -5.297
57071
   -11.99493122   -33.1705246    -40.28797531]
[ -33.35050201    -4.46615314    -8.78530598     7.71556234    -6.949
57018
   -7.67655945   -47.17803955   -16.71727753   -16.28227043     0.770
72573
   -31.82898712   -34.33365631   -12.14495373    12.36383152     4.254
0307
     7.12515688   -46.0978775    -14.60758114   -17.73284531   -24.223
29712
   -26.09648514   -52.14083862     1.0644542    -42.82314682   -19.951
75362
   -11.42463303    -7.85647058   -51.46764374    -9.00272751    -8.336
936
   -54.39671707   -37.38504791   -24.6989212     1.53812647   -19.004
77219
     5.35240459    -7.40578794   -17.2496357    -18.35800743    -6.729
59661
   -23.40646362   -36.33037186   -55.51404572]
[  -8.19537067     9.81828594   -25.48621559   -23.54010582    -4.855
68237
   -6.94281721   -17.68091583   -48.93915939   -34.76464844   -52.765
2092
   -50.41648865   -14.70374012     8.66214371   -27.15566635     6.447
21794
   -9.76159573   -59.79570007   -23.43365479    10.43781567   -56.436
71036
   -24.168396    -43.78590775   -52.15834808   -48.17084885   -38.039
3219
   -11.76746464   -14.47885036   -38.14334488   -47.70761871   -25.057
17087
   -35.11039734   -50.30174255    -8.42114544    -5.52239132   -41.710
50262
   -23.62919235   -21.55352402   -14.43875504     5.75120687   -34.407
27234
     22.58751869   -29.49789238   -41.74700546]
[ -71.62384796   -43.7921257    -51.80265045   -22.00750351  -103.656
52466

```

```

-17.76538467 -65.30267334 -43.8684082 -93.46403503 -34.294
49844
-9.96207142 7.04454041 -23.77471352 -10.41671562 -49.051
77307
-55.34400558 -59.59599686 -34.96619797 -25.52444267 -12.441
21361
-9.41991901 -9.95913887 6.64266682 2.65481162 -15.480
34191
25.82727432 -19.84401321 -48.59305573 -23.88368797 -1.995
40162
16.85153961 -22.45292664 -52.59412003 -24.38441467 15.534
77764
-14.75545406 -22.60277939 -34.56523895 -15.92673492 -75.492
72919
-36.34951782 -34.59121704 -60.93301773]
[ -34.50154877 -19.28138351 -18.63120842 -9.67574596 -48.081
60782
-22.70354652 -40.78736115 -21.65718269 -38.04601288 10.550
99487
-28.67878342 -15.69770241 16.99584007 4.14944553 -1.450
95742
-3.30947709 -39.31541061 -15.2105341 -28.85735512 -28.477
15378
-31.38887787 -49.80951309 -12.72123528 -19.36704826 -22.378
7117
-23.6171627 -13.96886444 -53.98962784 -3.78920889 -8.172
67036
-21.86774254 -58.35541534 -13.96012974 -8.60382366 8.489
84814
14.00068855 -8.69783878 -20.01979637 -0.48454282 -34.449
13483
-7.92018604 -15.24727726 -36.21956635]]
(5, 43)

```

softmax

```

[[ 1.72576594e-17 7.84864798e-02 7.32731714e-05 2.54908761e
-09
2.25045958e-12 3.24911234e-14 1.14750424e-18 1.42727201e
-15
9.31288195e-14 1.31959516e-19 3.00324433e-20 2.17065480e
-11
3.08444491e-04 1.15832000e-03 9.19946015e-01 4.16899425e
-13
1.27251267e-35 4.15206194e-12 3.77405403e-16 4.06054904e
-26
2.69678064e-18 1.70822124e-22 5.28217563e-18 9.65978667e
-26
4.87001974e-15 3.78089365e-11 5.17645427e-10 1.08103062e
-28
1.28490440e-17 9.50778403e-17 5.85471096e-19 1.40070859e
-21
7.49386427e-18 2.72885809e-05 3.30446434e-12 2.17986038e
-08

```

	1.55289271e-11	1.84979854e-07	7.19763875e-12	1.92972394e
-09				
	2.38159709e-12	1.51504703e-21	1.22844280e-24]	
[1.37901412e-20	4.82927121e-08	6.42831843e-10	9.42613091e
-03				
	4.03039024e-09	1.94814076e-09	1.36252795e-26	2.30826832e
-13				
	3.56620793e-13	9.08300080e-06	6.31472135e-20	5.15929350e
-21				
	2.23367956e-11	9.84126866e-01	2.95807724e-04	5.22304280e
-03				
	4.01285870e-26	1.90333578e-12	8.36047201e-14	1.26900517e
-16				
	1.94961210e-17	9.52860220e-29	1.21841103e-05	1.06084348e
-24				
	9.09013482e-15	4.59042179e-11	1.62736613e-09	1.86807368e
-28				
	5.17217047e-10	1.00651887e-09	9.98417054e-30	2.43998793e
-22				
	7.88683097e-17	1.95662360e-05	2.34336544e-14	8.87211296e
-04				
	2.55396326e-09	1.35545546e-13	4.47430416e-14	5.02204722e
-09				
	2.87216723e-16	7.00531506e-22	3.26635951e-30]	
[4.27717079e-14	2.84700809e-06	1.32384905e-21	9.26879383e
-21				
	1.20660535e-12	1.49669692e-13	3.24816373e-18	8.63681426e
-32				
	1.23670784e-25	1.88244069e-33	1.97132989e-32	6.37687166e
-17				
	8.95947323e-07	2.49347143e-22	9.78030670e-08	8.93211913e
-15				
	1.66502023e-36	1.03099161e-20	5.28989449e-06	4.78860009e
-35				
	4.94494104e-21	1.49410922e-29	3.45365861e-33	1.86220134e
-31				
	4.67836926e-27	1.20175499e-15	7.98503584e-17	4.21615185e
-27				
	2.95941964e-31	2.03315017e-21	8.75206045e-26	2.21102112e
-32				
	3.41275199e-14	6.19466175e-13	1.19047688e-28	8.47876526e
-21				
	6.75746654e-20	8.31168020e-17	4.87617058e-08	1.76796462e
-25				
	9.99990821e-01	2.39656606e-23	1.14780032e-28]	
[0.00000000e+00	5.81579498e-31	1.93054881e-34	1.68093248e
-21				
	0.00000000e+00	1.16917378e-19	0.00000000e+00	5.38867111e
-31				
	0.00000000e+00	7.75132954e-27	2.86296104e-16	6.96133684e
-09				
	2.87117652e-22	1.81704780e-16	3.02256001e-33	5.59359998e
-36				
	7.96296989e-38	3.95968654e-27	4.99072789e-23	2.39959554e

```

-17      4.92344624e-16      2.87136043e-16      4.65758454e-09      8.63490876e
-11      1.14884158e-18      9.99839664e-01      1.46265598e-20      4.78178502e
-33      2.57474683e-22      8.25459655e-13      1.26420739e-04      1.07672673e
-21      8.74886179e-35      1.56052553e-22      3.38809732e-05      2.37178304e
-18      9.26883745e-22      5.91282564e-27      7.35181481e-19      0.00000000e
+00      9.92870620e-28      5.76119987e-27      2.09129012e-38]
[      4.10227550e-23      1.67132253e-16      3.20204130e-16      2.48161994e
-12      5.19136824e-29      5.45547948e-18      7.64066637e-26      1.55332723e
-17      1.18490537e-24      1.51041662e-03      1.38618555e-20      6.01772524e
-15      9.50730622e-01      2.50575840e-06      9.26221322e-09      1.44401058e
-09      3.32960636e-25      9.79504665e-15      1.15949218e-20      1.69585225e
-20      9.22232975e-22      9.22278498e-30      1.18058056e-13      1.53410730e
-16      7.54931763e-18      2.18803603e-18      3.39044859e-14      1.41078647e
-31      8.93768892e-10      1.11560969e-11      1.25839789e-17      1.79234750e
-33      3.42018961e-14      7.24878447e-12      1.92287975e-04      4.75641415e
-02      6.59833082e-12      7.98676300e-17      2.43457769e-08      4.32302686e
-23      1.43603298e-11      9.44167636e-15      7.36036412e-24]]
(5, 43)
0 14 0.919946
1 15 0.00522304
2 27 4.21615e-27
3 30 0.000126421
4 40 1.43603e-11

```

Analyze Performance

```

In [17]: with tf.Session() as sess:
          saver.restore(sess, tf.train.latest_checkpoint('./lenet'))
          _, test_2_accuracy = evaluate(xtest21, ylabels)
          print("Test 2 Accuracy = {:.3f}".format(test_2_accuracy))

```

```

INFO:tensorflow:Restoring parameters from ./lenet/lenet
Test 2 Accuracy = 0.200

```

Output Top 5 Softmax Probabilities For Each Image Found on the Web

For each of these 5 new images, show the 5 highest probabilities.

```
In [18]: with tf.Session() as sess:
          saver.restore(sess, tf.train.latest_checkpoint('./lenet'))

          prob1 = sess.run(tf.nn.top_k(logits, k=5), feed_dict={x: xtest2
1})
          print(prob1)

          prob2 = sess.run(tf.nn.top_k(tf.nn.softmax(logits), k=5), feed_
dict={x: xtest21})
          print(prob2)
```

```
INFO:tensorflow:Restoring parameters from ./lenet/lenet
TopKV2(values=array([[ 14.6848774 ,  12.22348881,   8.00753307,
 6.68434906,
   5.24700212],
 [ 12.36383152,   7.71556234,   7.12515688,   5.35240459,
 4.2540307 ],
 [ 22.58751869,  10.43781567,   9.81828594,   8.66214371,
 6.44721794],
 [ 25.82727432,  16.85153961,  15.53477764,   7.04454041,
 6.64266682],
 [ 16.99584007,  14.00068855,  10.55099487,   8.48984814,
 4.14944553]], dtype=float32), indices=array([[14,  1, 13
, 12,  2],
 [13,  3, 15, 35, 14],
 [40, 18,  1, 12, 14],
 [25, 30, 34, 11, 22],
 [12, 35,  9, 34, 13]], dtype=int32))
TopKV2(values=array([[ 9.19946015e-01,   7.84864798e-02,   1.1583
2000e-03,
   3.08444491e-04,   7.32731714e-05],
 [ 9.84126866e-01,   9.42613091e-03,   5.22304280e-03,
 8.87211296e-04,   2.95807724e-04],
 [ 9.99990821e-01,   5.28989449e-06,   2.84700809e-06,
 8.95947323e-07,   9.78030670e-08],
 [ 9.99839664e-01,   1.26420739e-04,   3.38809732e-05,
 6.96133684e-09,   4.65758454e-09],
 [ 9.50730622e-01,   4.75641415e-02,   1.51041662e-03,
 1.92287975e-04,   2.50575840e-06]], dtype=float32), indi
ces=array([[14,  1, 13, 12,  2],
 [13,  3, 15, 35, 14],
 [40, 18,  1, 12, 14],
 [25, 30, 34, 11, 22],
 [12, 35,  9, 34, 13]], dtype=int32))
```

Model was able to predict only 1 out of 5 images, the first image with 91%. This image contains the "stop" sign, which is the 14th label.

In []: